

AI in Marketing, Personalisation and Social Media

WHY THIS MODULE MATTERS

Personalisation and targeted marketing are core to modern retail, but the line between helpful and manipulative is thinner than it looks. This module gives you a framework for telling the difference.

6.1 Where Helpful Becomes Manipulative

Personalisation is central to modern retail marketing: recommending products a customer is likely to want, timing an email for when they usually shop, tailoring a homepage to past browsing. Done well, this is genuinely useful — customers find what they wanted faster. The ethical question is not whether to personalise, but where personalisation stops serving the customer's interest and starts working against it.

A useful test: would the customer be comfortable if you explained exactly how and why they were shown this, at this moment, at this price? If the honest answer is no, the technique is probably relying on the customer not knowing.

6.2 Practices That Cross the Line

- **Exploiting inferred vulnerability** — targeting based on inferred financial distress, body image insecurity, or life events the customer never disclosed. Fashion is particularly exposed here, since body-related insecurity is easy to infer and easy to monetise.
- **Dark patterns and false urgency** — AI-generated “only 2 left” or “14 people viewing” messages that are not tied to real stock or real traffic. UK consumer protection law treats materially misleading urgency claims as unfair commercial practices.
- **Opaque personalised pricing** — showing different prices to different individuals based on inferred willingness to pay, without disclosure. Legally fraught and, once discovered, reliably damaging to trust.
- **Undisclosed AI-generated influencer or review content** — synthetic testimonials, AI-written reviews, or virtual influencers presented as real people without disclosure.

6.3 A Decision Framework for Personalisation Features




Figure 6.1 — A decision path for assessing a proposed personalisation feature before launch.

Three questions resolve most cases. First: does this use personal or inferred behavioural data? If so, the Module 5 obligations apply. Second: could it influence price, access or opportunity for an individual? If so, it needs a fairness review and probably a DPIA. Third: would the customer feel misled if the mechanism were explained to them? If so, either disclose it or do not ship it.

6.4 Social Media, Influencers and Disclosure

The Advertising Standards Authority and the Competition and Markets Authority both require that advertising is clearly identifiable as such. AI adds two newer wrinkles: content that is wholly AI-generated but presented as authentic creator output, and virtual influencers that some audiences may not realise are synthetic. The safe position is the same in both cases — disclose prominently, in the content itself rather than buried in a profile bio.

TRY THIS — YOU ARE THE MARKETING GOVERNANCE LEAD

Scenario: your growth team proposes a campaign that uses browsing and purchase data to identify customers whose behaviour suggests they are dieting, then targets them with smaller-size product recommendations and time-limited offers. Projected conversion uplift is significant. What is your call? (Model answer: reject the targeting as designed. Inferring dieting behaviour edges into health-related inference, which brings Article 9 special category considerations into play, and targeting on inferred body insecurity fails the ‘would the customer be comfortable if we explained it’ test regardless of the legal position. A defensible alternative is letting customers explicitly set size and fit preferences themselves — same commercial benefit, customer-declared rather than covertly inferred.)

IN PRACTICE

A retailer audits its urgency messaging and finds a third of “low stock” banners fire on a fixed threshold rather than real inventory. The team ties every banner to live stock data and drops the ones that cannot be substantiated. Conversion falls slightly; complaint volume and return rates fall further.

KEY TAKEAWAYS

- The test for personalisation is whether the customer would be comfortable if the mechanism were fully explained to them.
- Targeting on inferred vulnerability — financial distress, body insecurity, life events — is the clearest line in retail marketing, and fashion is unusually exposed to it.
- False urgency and unsubstantiated scarcity claims are unfair commercial practices under UK consumer law, whether AI-generated or not.
- AI-generated influencer and review content requires prominent disclosure in the content itself, not buried in a bio.
- Customer-declared preferences achieve most of the commercial benefit of covert inference, without the ethical and legal exposure.

FURTHER READING

- Advertising Standards Authority — guidance on recognising advertising online, asa.org.uk
- Competition and Markets Authority — guidance on online choice architecture
- ICO — guidance on direct marketing and profiling

MODULE 6 CHECK-IN QUIZ

Five questions to confirm your understanding before moving on. Answers are provided in the Answer Key at the end of the course.

1. What is the practical test offered in Module 6 for whether personalisation has crossed a line?

- A. Whether it increases conversion
- B. Whether a competitor is doing the same thing
- C. Whether the customer would be comfortable if the mechanism were fully explained to them
- D. Whether it uses machine learning

2. Why is fashion retail particularly exposed to targeting on inferred vulnerability?

- A. Fashion customers are less price-sensitive
- B. Fashion is exempt from consumer protection law
- C. Fashion retailers rarely hold behavioural data
- D. Body-related insecurity is comparatively easy to infer and easy to monetise

3. How does UK consumer law treat AI-generated “only 2 left” messaging not tied to real stock?

- A. As a potentially unfair commercial practice through materially misleading urgency
- B. As acceptable marketing practice
- C. As exempt because it is automated
- D. As a matter solely for the ASA, with no legal dimension

4. Where should disclosure of AI-generated influencer content appear?

- A. In the profile bio only
- B. Prominently in the content itself
- C. In the brand's annual report
- D. Disclosure is not required

5. In the dieting-targeting exercise, what is the recommended alternative?

- A. Proceed as designed, since conversion uplift is significant
- B. Target only customers over 25
- C. Let customers explicitly set their own size and fit preferences instead of covertly inferring them
- D. Run the campaign without disclosure