

AI in Recruitment and Talent Acquisition

WHY THIS MODULE MATTERS

Algorithmic hiring tools can quietly discriminate at scale. This module gives recruitment and talent teams the tools to catch that before it happens, not after.

8.1 Hiring Is Where AI Bias Becomes a Legal Problem

Fashion and retail businesses recruit at volume — seasonal store staff, warehouse and distribution teams, head-office graduate intakes — which makes automated CV screening and candidate ranking commercially attractive. It is also the area where AI bias moves fastest from an ethical concern to a legal one, because hiring decisions are directly governed by the Equality Act 2010.

CASE STUDY — AMAZON'S ABANDONED RECRUITING TOOL

Between 2014 and 2017, Amazon built an internal AI tool to rank job applicants' CVs, trained on ten years of the company's own hiring data. Because that historical data came from a workforce that was predominantly male, particularly in technical roles, the model learned that male candidates were preferable. Reuters reported in 2018 that it penalised CVs containing the word “women's” — as in “women's chess club captain” — and downgraded graduates of two all-women's colleges.


Engineers edited the model to stop it penalising those specific terms, but could not be confident it had not learned other, subtler proxies for gender from the same data. Amazon scrapped the tool rather than deploy it. Nobody at Amazon set out to build a discriminatory hiring system: the bias was inherited wholesale from historical outcomes the model was trained to reproduce. Any retailer training a screening model on its own past hiring decisions is running the same risk.

8.2 The Equality Act 2010 and Indirect Discrimination

The Equality Act prohibits discrimination on the basis of protected characteristics including age, disability, race, sex, religion or belief, and others. Critically, **indirect** discrimination does not require intent: a screening practice applied equally to everyone can still be unlawful if it puts people sharing a protected characteristic at a particular disadvantage and the employer cannot objectively justify it. An AI model that systematically ranks one group lower falls squarely within this, whether or not anyone intended it and whether or not the protected characteristic was ever an input.

Disability discrimination deserves specific attention in AI screening. Video-interview analysis tools that score facial expression, speech patterns or response timing can systematically disadvantage neurodivergent candidates or those with speech, hearing or motor conditions — and employers also carry a duty to make reasonable adjustments, which is difficult to honour if the first sift is fully automated.

8.3 Governing Recruitment AI



Bias testing and human review are not optional extras

Figure 8.1 — A decision path for approving an AI screening or ranking tool.

- **Keep a human in the loop before rejection** — no candidate should be rejected solely by an automated score. This is both an Article 22 safeguard and the single most effective bias control.
- **Test for adverse impact before launch and on a schedule after** — compare selection rates across protected groups, not just overall model accuracy.
- **Interrogate the training data** — if a model is trained on past hiring decisions, it will reproduce whatever bias those decisions contained.
- **Offer adjustments and an alternative route** — candidates must be able to request a non-automated process without disadvantage.
- **Tell candidates AI is being used** — transparency is a UK GDPR requirement and, increasingly, a candidate expectation.
- **Interrogate vendors** — ask what data a third-party tool was trained on, what adverse-impact testing exists, and what contractual audit rights you hold.

TRY THIS — YOU ARE THE TALENT ACQUISITION LEAD

Scenario: a vendor offers a CV-screening tool promising to cut time-to-shortlist by 70%, trained on “millions of successful hires across the retail sector”. It auto-rejects the bottom 40% of applicants before any human sees them. What do you require before adopting it? (Model answer: the auto-rejection is the dealbreaker as configured — rejection with no human review is exactly the solely-automated decision Article 22 restricts, and removes the most effective bias control you have. Before adoption: require adverse-impact testing across protected characteristics on your own applicant pool, not just the vendor's claims; clarify what ‘successful hires’ meant in the training data, since it likely encodes the sector's existing demographic skew; reconfigure so the tool ranks and assists rather than auto-rejects; and secure contractual audit rights plus a candidate-facing disclosure and adjustment route.)

IN PRACTICE

A retailer piloting a screening tool for seasonal hiring runs adverse-impact testing before go-live and finds the model down-ranks applicants with employment gaps — a pattern that disproportionately affected women returning from maternity leave and candidates with disabilities. The feature is removed from scoring before a single candidate is affected.

KEY TAKEAWAYS

- Recruitment is where AI bias moves fastest from an ethical concern to a legal one, under the Equality Act 2010.
- Amazon's abandoned tool shows how a model trained on past hiring decisions inherits their bias wholesale, without anyone intending it.
- Indirect discrimination requires no intent — a uniformly applied model that disadvantages a protected group can be unlawful unless objectively justified.
- A human reviewing before any rejection is both an Article 22 safeguard and the single most effective bias control available.
- Adverse-impact testing must run on your own applicant pool, before launch and on a schedule afterwards — vendor assurances are not a substitute.

FURTHER READING

- Reuters — “Amazon scraps secret AI recruiting tool that showed bias against women” (2018)
- Equality and Human Rights Commission — guidance on the Equality Act 2010
- ICO — guidance on AI and data protection in recruitment

MODULE 8 CHECK-IN QUIZ

Five questions to confirm your understanding before moving on. Answers are provided in the Answer Key at the end of the course.

1. What happened with Amazon's AI recruiting tool?

- A. Trained on ten years of predominantly male hiring data, it learned to penalise CVs indicating female candidates, and was scrapped
- B. It was deployed successfully across the company
- C. It was blocked by regulators before completion
- D. It only affected warehouse hiring

2. What is required for indirect discrimination under the Equality Act 2010?

- A. Proof that the employer intended to discriminate
- B. No intent is required — a uniformly applied practice can be unlawful if it disadvantages a protected group without objective justification
- C. A protected characteristic must be an explicit model input
- D. A complaint from at least ten candidates

3. Why do video-interview analysis tools warrant particular caution?

- A. They are slower than human review
- B. They cannot process CVs
- C. Scoring facial expression, speech patterns or timing can disadvantage neurodivergent and disabled candidates, and complicate reasonable adjustments
- D. They are prohibited outright in the UK

4. What is described as both an Article 22 safeguard and the single most effective bias control in recruitment AI?

- A. Using a larger training dataset
- B. Removing all demographic data from CVs
- C. Running the tool only for senior roles
- D. A human reviewing before any candidate is rejected

5. In the vendor exercise, what is identified as the dealbreaker in the tool as configured?

- A. Auto-rejecting the bottom 40% of applicants with no human review
- B. The price of the licence
- C. The vendor being based overseas
- D. The 70% time-to-shortlist claim